

Application of integer linear programming to scheduling clinician on-call service in Toronto hospital: case study

David Landsman^{a, b}, Huiting Ma^{*a}, Jesse Knight^{*a, c}, Kevin Gough^d, and Sharmistha Mishra^{†a, c, d, e}

^a*MAP Centre for Urban Health Solutions, St. Michael's Hospital, Unity Health Toronto, Toronto, ON, Canada*

^b*Department of Computer Science, University of Toronto, Toronto, ON, Canada*

^c*Institute of Medical Sciences, University of Toronto, Toronto, ON, Canada*

^d*Department of Medicine, Division of Infectious Disease, St. Michael's Hospital, Unity Health Toronto, Toronto, ON, Canada*

^e*Institute of Health Policy, Management and Evaluation, Dalla Lana School of Public Health, University of Toronto, Toronto, ON, Canada*

Abstract

Scheduling of personnel in a hospital environment is vital to improving the service provided to patients and balancing the workload assigned to clinicians. Many approaches have been tried and successfully applied to generate efficient schedules in such settings. However, due to the computational complexity of the scheduling problem in general, most approaches resort to heuristics to find a non-optimal solution in a reasonable amount of time. We designed an integer linear programming formulation to find an optimal schedule in a clinical division of a hospital. Our formulation mitigates issues related to computational complexity by minimizing the set of constraints, yet retains sufficient flexibility so that it can be adapted to a variety of clinical divisions.

We then conducted a case study for our approach using data from the Infectious Diseases division at St. Michael's Hospital in Toronto, Canada. We analyzed and compared the results of our approach to manually-created schedules at the hospital, and found improved adherence to departmental constraints and clinician preferences. We used simulated data to examine the sensitivity of the runtime of our linear program for various parameters and observed reassuring results, signifying the practicality and generalizability of our approach in different real-world scenarios.

Keywords: Clinician scheduling, Integer linear programming, Exact method, Optimization, Clinician requests

1 Introduction

Hospital departments must allocate and use their limited resources efficiently in order to provide a high quality of care for their patients. In particular, on-call schedules for a fixed number of health-care providers are central

^{*}Contributed equally to this work.

[†]Corresponding author (sharmistha.mishra@utoronto.ca)

Preprint submitted to arXiv: <https://arxiv.org/abs/1910.08526>

to the efficient running of hospitals. Carefully allocated schedules should balance sufficient staff with workload to maximize quality of care. It is common for on-call schedules in hospitals to be created manually. However, manually-created schedules are subject to three problems. First, when there are a large number of clinicians, or the constraints that need to be satisfied by the schedule are complex, it becomes infeasible to find a satisfying schedule by hand. Second, when creating a schedule manually it is difficult to ensure that all constraints are met while also trying to satisfy all staff preferences. Third, manual scheduling is often time-consuming even for relatively small departments, and can take up time and resources that are better used for improving patient care. For these reasons, it is important to develop automated methods that can efficiently generate schedules that satisfy the given constraints.

Automated methods to generate schedules have been studied and applied in many industries, including transportation [1, 2, 3], retail [4, 5], and military [6, 7]. Of special interest to a clinician scheduling problem are the approaches to scheduling nurses, who often work in shifts. In the nurse scheduling problem, the goal is to find an optimal assignment of nurses to shifts that satisfies all hard constraints (such as hospital regulations), and as many soft constraints (such as nurse preferences) as possible. Hard constraints must be satisfied by any candidate solution to the nurse scheduling problem, while soft constraints can be used to rank the candidate solutions. For instance, a nurse scheduling problem may include a hard constraint to assign at most a single shift for each nurse per day. It can also incorporate nurse preference for shift time (that is, day versus night shifts) as a soft constraint that is meant to optimize the schedule, but is not guaranteed to be fulfilled. A wide variety of approaches, including exact and heuristic approaches, have been used to solve the nurse scheduling problem: integer linear programming [8, 9, 10], network flows [11], genetic algorithms [12, 13, 14], simulated annealing [15], and artificial intelligence [16, 17]. A comprehensive literature review of these and other methods applied to nurse scheduling is presented in [18].

Many of the approaches to nurse scheduling were designed to satisfy the requirements of a specific hospital department, which results in a large number of variables and constraints to be incorporated into the problem formulation. While these department-specific approaches allow end-users to find precise schedules that satisfy the needs of that department and the preferences of nurses and clinicians in that department, they are difficult to adapt to other departments. Moreover, the large number of variables and constraints also leads to computational complexity issues [19], especially when trying to find the most optimal solution. In particular, difficult instances of these formulations become impossible to solve in a reasonable amount of time. In this paper, we tackle the clinician scheduling problem arising from a case study of one clinical division, providing two different services (general infectious disease (ID) consults; and HIV consults) at St. Michael’s Hospital in Toronto, Canada. The clinician scheduling problem involves creating a yearly schedule that assigns clinicians to on-call work on a weekly basis, while ensuring a fair and balanced workload. Our goals in this paper are to (1) present an integer linear programming (ILP) formulation for our problem, and describe the flexibility of this formulation for solving similar problems; (2) compare the performance of our tool for solving the ILP formulation to the results of a manual approach; and (3) analyze the robustness of this approach in difficult instances of the problem, by exploring the change in runtime with changes to: the number of clinicians, the number of services provided by the department, the number of requests per clinician, and the time-horizon of the schedule.

We begin by describing the details of the problem in Section 2, and presenting our ILP formulation in Section 3. We then categorize our nurse rostering problem and solution approach in the categorization scheme of [20] and compare our problem to recent studies in Section 4. Next, we analyze the effectiveness of our automated approach when compared to manually-created schedules, and evaluate the performance of the algorithm on simulated data in Section 5. Finally, we discuss and interpret the results in Section 7.

2 Problem Description

At St. Michael’s Hospital, the division of infectious diseases offers separate but concurrent services for general ID consultations and for HIV consultations. Each service provides clinical care throughout the year, during regular working hours and on weekends and holidays. The schedule is created in advance, outlining all work-week and weekend shifts for the full year. In the yearly schedule clinicians are assigned to “blocks” of two

consecutive regular work weeks and individual weekends. Apart from long (holiday) weekends, a work week starts on Monday at 8 A.M. and ends on Friday at 5 P.M. Accordingly, weekend service starts on Friday at 5 P.M., and ends on Monday at 8 A.M. During the weekend, ID and HIV consultation services are combined and provided by one clinician. During the regular work week the ID and HIV services are led by one clinician each. Therefore, our objective is to assign a single clinician to cover each service for each block and each weekend of the year, while additionally ensuring a balanced workload.

In most scheduling problems, the constraints can be divided into hard and soft constraints. Hard constraints must be satisfied by any candidate solution, while soft constraints can be used to select a more favourable solution from a set of candidate solutions. Typically, soft constraints are therefore encoded as objective functions which are maximized when finding a solution. In the case of the clinician scheduling problem, we chose departmental regulations and workload balance as hard constraints, while clinician preference and requests serve as soft constraints. After the schedule is generated, clinicians may exchange certain weeks or days throughout the year, to fulfill any missed requests.

The following are the departmental and workload constraints placed on the clinician assignments. First, each clinician must work between a minimum and maximum number of blocks of each service during the year. For instance, one clinician might have to provide 3-5 blocks of general ID service and 2-3 blocks of HIV service throughout the year. These limits may be different for each clinician, and they may change from year to year as the number of clinicians in the department changes. Second, the schedule should not assign a clinician to work for two consecutive blocks or two consecutive weekends. The schedule should also distribute regular weekends and holiday weekends each equally among all clinicians.

In addition to balancing the workload among clinicians, the schedule should accommodate their preferences. Clinicians provide their requests for time off before schedule generation so that the requests may be integrated into the schedule. Clinicians may specify days, weeks or weekends off, with the understanding that any blocks overlapping with their request will be assigned to a different clinician where possible. For example, if a clinician only requests a given Monday and Tuesday off, the schedule should avoid assigning the entire block to that clinician. Clinicians also typically prefer to have their weekend and block assignments side by side, so the schedule should accommodate this where possible. A summary of the outlined constraints is given in Table 1.

3 Mathematical Formulation

In this section, we present an application of integer linear programming to solve the clinician scheduling problem presented in Section 2. An integer linear program (ILP) consists of a linear objective function and linear constraints, with integer-valued variables. The optimal solution to the ILP must lie within the space defined by the constraints while also maximizing the objective value. First, we describe the sets, indices and variables necessary for the formulation of the problem. We then write the constraints given in Table 1 as linear functions of the variables, and define the objective function of the ILP.

3.1 Sets and Indices

We denote the set of all services that clinicians in a single department can provide as \mathcal{S} . The set of all clinicians in the department is denoted as \mathcal{C} . The sets of blocks and weekends are denoted as \mathcal{B} and \mathcal{W} respectively. The block size used in our experiments is 2 weeks, but the following LP formulation does not require a particular size for blocks, and so it can be adapted for other cases. A subset of weekends are denoted $\mathcal{L} \subset \mathcal{W}$, corresponding to statutory long/holiday weekends such as the Thanksgiving weekend in the United States and Canada. Lastly, the block and weekend time-off requests of clinicians are denoted as the following subsets of all blocks and weekends: $\mathcal{U}_c \subset \mathcal{B}$ and $\mathcal{V}_c \subset \mathcal{W}$, respectively. For instance, if clinician c 's time-off requests intersect with blocks 1 and 2, and weekend 1, then $\mathcal{U}_c = \{1, 2\}$ and $\mathcal{V}_c = \{1\}$. Table 2 presents a summary of the sets and indices described.

Table 1: Summary of the constraints for the clinician scheduling problem

Constraint Name	Abbreviation	Description	Type
Block Coverage	BC	each service needs to have exactly one clinician that covers any given block	Hard
Weekend Coverage	WC	every weekend needs to have exactly one clinician that covers it	Hard
Min/Max	MM	for a given service, each clinician can only work between the minimum and maximum number of allowed blocks	Hard
No Consecutive Blocks	NCB	any clinician should not work two consecutive blocks, across all services	Hard
No Consecutive Weekends	NCW	any clinician should not work two consecutive weekends	Hard
Equal Weekends	EW	weekends should be equally distributed between clinicians	Hard
Equal Holidays	EH	long weekends should be equally distributed between clinicians	Hard
Block Requests	BR	each clinician can request to be off service during certain blocks throughout the year	Soft
Weekend Requests	WR	each clinician can request to be off service during certain weekends throughout the year	Soft
Block-Weekend Adjacency	BWA	the block and weekend assignments of a given clinician should be adjacent	Soft

Hard constraints must be satisfied by any candidate schedule. Soft constraints are optionally satisfied, and are used to rank the set of candidate solutions.

3.2 Variables

Since each clinician may be assigned to work on any service, during any block of the year, we denote such an assignment as a binary variable $X_{c,b,s} \in \{0,1\}$. A value of 1 indicates that clinician c is assigned to service s during block b , while a value of 0 indicates they are not assigned. Weekend assignments are similarly defined using a binary variable $Y_{c,w} \in \{0,1\}$, but without a service index, as clinicians are expected to provide all services during the weekends. We then define $m_{c,s}$ and $M_{c,s}$ to represent the minimal and maximal number of blocks of service s that clinician c is required to work during the year. Table 3 presents a summary of the constants and variables in the problem.

3.3 Constraints

We now formalize the hard constraints in Table 1 using the variables defined above. The BC (block coverage) and WC (weekend coverage) constraints, are given by Eqns. (1) and (2), respectively.

$$\sum_{c=1}^C X_{c,b,s} = 1 \quad \forall b \in \mathcal{B}, s \in \mathcal{S} \quad (1)$$

$$\sum_{c=1}^C Y_{c,w} = 1 \quad \forall w \in \mathcal{W} \quad (2)$$

Table 2: Description of sets and indices in the problem

Set	Index	Description
$\mathcal{S} = \{1, \dots, S\}$	s	services
$\mathcal{C} = \{1, \dots, C\}$	c	clinicians
$\mathcal{B} = \{1, \dots, B\}$	b	blocks
$\mathcal{W} = \{1, \dots, W\}$	w	weekends
$\mathcal{L} \subset \mathcal{W}$		long weekends
$\mathcal{U}_c \subset \mathcal{B}$		block requests of clinician c
$\mathcal{V}_c \subset \mathcal{W}$		weekend requests of clinician c

Table 3: Description of variables and constants in the problem

Name	Description
$X_{c,b,s} \in \{0, 1\}$	assignment of clinician c to service s on block b
$Y_{c,w} \in \{0, 1\}$	assignment of clinician c on weekend w
$m_{c,s}$	minimum number of blocks clinician c should cover on service s
$M_{c,s}$	maximum number of blocks clinician c should cover on service s

The MM (min/max) constraint is given by:

$$m_{c,s} \leq \sum_{b=1}^B X_{c,b,s} \leq M_{c,s} \quad \forall c \in \mathcal{C}, s \in \mathcal{S} \quad (3)$$

The NCB (no consecutive blocks) and NCW (no consecutive weekends) constraints are given by Eqns. (4) and (5), respectively.

$$\sum_{s=1}^S (X_{c,b,s} + X_{c,b+1,s}) \leq 1 \quad \forall c \in \mathcal{C}, b \leq B-1 \quad (4)$$

$$Y_{c,w} + Y_{c,w+1} \leq 1 \quad \forall c \in \mathcal{C}, w \leq W-1 \quad (5)$$

The EW (equal weekend) and EH (equal holidays) constraints are given by Eqns. (6) and (7), respectively.

$$\left\lfloor \frac{W}{C} \right\rfloor \leq \sum_{w=1}^W Y_{c,w} \leq \left\lceil \frac{W}{C} \right\rceil \quad \forall c \in \mathcal{C} \quad (6)$$

$$\left\lfloor \frac{|\mathcal{L}|}{C} \right\rfloor \leq \sum_{w \in \mathcal{L}} Y_{c,w} \leq \left\lceil \frac{|\mathcal{L}|}{C} \right\rceil \quad \forall c \in \mathcal{C} \quad (7)$$

3.4 Objectives

As described in Section 2, the soft constraints of the clinician scheduling problem include: satisfying clinician block off requests (BR), satisfying clinician weekend off requests (WR), and assigning weekends closer to blocks (BWA). We convert these soft constraints into linear objective functions of the binary variables defined in Section 3.2. Objectives BR and WR are given in Eqns. (8) and (9) as linear functions of X and Y :

$$Q_1(X) = \sum_{c=1}^C \sum_{b=1}^B \sum_{s=1}^S (-1)^{\mathbf{1}(b \in \mathcal{U}_c)} \cdot X_{c,b,s} \quad (8)$$

$$Q_2(Y) = \sum_{c=1}^C \sum_{w=1}^W (-1)^{\mathbf{1}(w \in \mathcal{V}_c)} \cdot Y_{c,w} \quad (9)$$

where $\mathbf{1}(P)$ is the indicator function that has value 1 when predicate P holds and 0 otherwise. In the above two objectives, we penalize any assignments that conflict with a block or weekend request, and aim to maximize the non-conflicting assignments.

The BWA objective is optimized by considering the product $X_{c,b,s} \cdot Y_{c,w}$ for values of w “adjacent” to the value of b . This leads to the maximization objective:

$$Q_3(X, Y) = \sum_{c=1}^C \sum_{b=1}^B \sum_{s=1}^S X_{c,b,s} \cdot Y_{c,w=\varphi(b)} \quad (10)$$

where $\varphi(b)$ is a one-to-one mapping of a block to an adjacent weekend, by some appropriate definition of adjacency. For instance, clinicians might want to be assigned during a weekend that falls within an assigned block. In this case, we will have $\varphi(b) = 2b - 1$.

However, as it is, Q_3 is not a linear function of the assignment variables X and Y , and cannot be optimized in a linear programming framework. An approach used to convert such functions into linear objectives involves introducing a helper variable and additional constraints [21]. In our case, we introduce a variable $Z_{c,b,s}$ for every product $X_{c,b,s} \cdot Y_{c,w}$ with $w = \varphi(b)$, and constraining Z such that

$$Z_{c,b,s} \leq X_{c,b,s} \quad (11)$$

$$Z_{c,b,s} \leq Y_{c,w=\varphi(b)} \quad \forall s \in \mathcal{S} \quad (12)$$

Since $X_{c,b,s}$ and $Y_{c,w}$ are binary variables, $Z_{c,b,s}$ will be constrained to 0, unless both $X_{c,b,s}$ and $Y_{c,w}$ are 1. Therefore, it suffices to maximize the following linear function of Z ,

$$Q_3(Z) = \sum_{c=1}^C \sum_{b=1}^B \sum_{s=1}^S Z_{c,b,s} \quad (13)$$

to get the correct adjacency maximization objective.

In order to optimize all objectives simultaneously, we optimize a weighted sum of the normalized objective functions,

$$\max_{X,Y,Z} \alpha_1 \bar{Q}_1(X) + \alpha_2 \bar{Q}_2(Y) + \alpha_3 \bar{Q}_3(Z) \quad (14)$$

subject to the constraints defined in Section 3.3, where \bar{Q}_i is the normalization of objective Q_i , and $0 \leq \alpha_i \leq 1$. This method guarantees an optimal solution to be Pareto optimal [22].

One of the efficient approaches to finding an exact solution for an ILP is Branch-and-Cut [23]. This method involves iteratively solving LP relaxations of the ILP, then constraining the relaxed problems and considering various sub-problems until it finds integral solutions to the original ILP. In the intermediate relaxations, the integer assignment variables can take on real values, allowing the problem to be solved efficiently using the Simplex method [24]. The complexity of finding an optimal integral solution thus lies in the branching search structure of Branch-and-Cut.

4 Categorization and Comparison with Other Models

In 2010, De Causmaecker and Berghe [20] addressed a gap in the literature surrounding a central categorization system for the nurse rostering problems identified in the literature. The authors created a categorization framework that takes into account hard and soft constraints related to nurse preferences, consecutive assignments, coverage and various flexibilities in terms of shift and skill types. The described framework is inspired by the categorization used in theoretic scheduling problems which focuses on three main aspects of scheduling: (1) the machine environment; (2) the job characteristics; and (3) the objective function. De Causmaecker and Berghe use these three aspects as a basis for their categorization, and adjust them to fit the realm of nurse rostering. A description of the categories defined by the authors is given in Table 4.

Table 4: Description of De Causmaecker and Berghe’s categorization system for nurse rostering problems.

Category	Sub-categories	
α : Personnel environment	Personnel Constraints	Skill Interactions
	(A) Availability	Fixed number
	(S) Sequences	Variable number
	(B) Balance (C) Chaperoning	Individual skill definitions
β : Work characteristics	Coverage Constraints	Shift Type
	(R) Range	Fixed number
	(T) Time intervals	Variable number
	(V) Fluctuating	Overlapping
γ : Optimization objective	Objective	Mode
	(P) Personnel constraints	Multi-objective
	(L) Coverage constraints	
	(X) Number of personnel	
	(R) Robustness	
	(G) General	

We follow the categorization in Table 4 to situate our clinician scheduling problem. Our problem includes personnel environment constraints targeting the availability and preferences of clinicians (constraints BR and WR in Table 1), sequences of assignments (constraints NCB and NCW), and balance of rosters (constraints EW and EH) which fall under the categories $\alpha : A$, $\alpha : S$ and $\alpha : B$, respectively. For work characteristics, the relevant category for our problem are the range constraints, $\beta : R$, as we restrict the workload of each clinician between a minimum and maximum allowed blocks (constraint MM). Lastly, for the optimization objective our problem corresponds to the category $\gamma : P$, as our optimization objective aims to adhere to clinician preferences. Therefore the appropriate categorization for our problem is $ASB/R/P$.

De Causmaecker and Berghe also present examples from the nurse rostering literature and their appropriate categorization in their framework. We use some of these examples to compare our clinician scheduling problem and solution approach with existing literature. Azaiez and Al Sharif [8] describe a nurse rostering problem to schedule nurses on a per-day basis, for a total period of 28 days. Azaiez and Al Sharif’s problem is categorized as $ASBN/V2/PL$. Similar to our clinician scheduling problem, the authors include hard constraints to prevent consecutive assignments, and enforce a minimum and maximum number of shifts for each nurse. Unlike our problem, Azaiez and Al Sharif choose an optimization goal that focuses on balancing the workload among nurses, preventing day-night consecutive shifts and avoiding off-on-off assignments. The authors do not attempt to incorporate nurse preferences for time off into the model.

In [25], Bilgin et al address a nurse rostering problem using local search. De Causmaecker and Berghe categorize the problem in [25] as $ASBI/RVNO/PLR$. Bilgin et al incorporate coverage constraints, overlapping assignment constraints and horizontal constraints, which include constraints arising from employment contracts (such as minimum and maximum thresholds). Unlike our approach for these constraints, the authors consider coverage constraints to be soft constraints, in order to prevent over-constraining of the problem. Bilgin et al also incorporate nurse requests into the model formulation by allowing nurses to request days off and days on.

Burke et al [26] present a flexible framework for nurse rostering based on various approaches observed from Belgin hospitals. In [20], Burke et al’s work is categorized as $ASBCI/RVNT0/PR$. The authors use a two-step algorithm comprised of an initialization procedure which provides a solution satisfying only the hard constraints of their problem, and a metaheuristic procedure, a hybrid tabu search, which aim to satisfy as many soft constraints as possible. Unlike our clinician scheduling problem, Burke et al’s hard constraints address only coverage requirements in an effort to make the first step of their algorithm efficient. In their metaheuristic procedure, the authors optimize soft constraints including assignment of free weekends and minimum number of assignments per staff member. Notably, Burke et al’s approach make use of “time interval personnel

requirements”, which describe the amount of required staff per skill category as a varying number throughout the work day.

5 Computational Experiments

Our goal was to determine if the schedules provided by solving the ILP could successfully (i) enforce all hard constraints; (ii) improve fulfillment of soft constraints compared to the manual approach; and (iii) assess whether our ILP formulation can be used for a wide range of configurations. First, we compared the schedules created by solving the ILP formulation given in Section 3 to schedules that were manually generated, with respect to adherence to the hard and soft constraints outlined in Section 2. We then examined the efficiency of the ILP approach in generating schedules by its runtime on a variety of instances that may be found in the real-world.

5.1 Implementation

We developed a Python software package with a user interface that implements the above linear program and allows configuration of clinicians, to be used by the ID division at St. Michael’s Hospital [27]. The software was used to generate the results in the following sections, using real data as well as simulated data as input. All the following experiments were conducted on an Intel Core i7-4770k CPU @ 3.50 GHz with 16 GB of RAM running 64-bit Windows 10. Our software package uses COIN-OR Branch-and-Cut open source solver version 2.9.9 [28].

5.2 Comparison with Manually Generated Schedules

We used clinician time-off requests and minimum/maximum requirements from 2015-2018 as input data for the ILP problem. Table 5 compares the optimal schedule generated using the software with the manually-created schedule for data from 2018. The schedule is color-coded to distinguish between the different clinicians.

First, we evaluated the ILP solution by comparing it with the manual generation, as in Table 5. Specifically, we examined the adherence of each schedule to the constraints presented in Table 1. As shown in Table 6, the ILP solution satisfied all hard constraints. In contrast, manual generation did not satisfy all hard constraints. In particular, we see that the manual generation assigned clinicians to multiple consecutive blocks in all four years. Moreover, the manual generation did not have an equal distribution of weekends and holidays for all four years of data. Considering all objectives, we see that the ILP solution outperforms manual generation in all four years, by accommodating almost all time-off requests and ensuring that weekends are always assigned close to blocks.

5.3 Influence of Problem Complexity on Runtime

Next, we examined the influence of the following four parameters on the runtime of the ILP solver using simulated data: number of clinicians; number of services offered; number of time-off requests per clinician per year; time-horizon of the schedule.

The effect of increasing the number of clinicians and number of services on the runtime of the program is shown in Table 7. We executed the algorithm for $S = \{1, 2, 3\}$ total services and $C = \{10, 20, 30, 50\}$ clinicians in total across all services. In a department providing a single service, increasing the number of clinicians did not affect the runtime, and we were able to find an ILP solution in all four cases within 1 second. For 2 concurrent services, a roster of 30 or more clinicians becomes impractical to schedule, as searching for a solution required over 24 hours. We saw similar issues for a roster of 20 or more clinicians assigned to a division with

Table 5: Comparison of automatically generated (ILP solution) and manually generated schedules for 2018. In the ILP solution there are rigid 2-week block assignments, unlike manual generation that often assigns 3-4 weeks in a row to a single clinician. Moreover, in the ILP solution we see for each block either the HIV or ID clinician was assigned weekend coverage, indicating improved Block-Weekend Adjacency.

Week #	ILP Solution			Manual Generation		
	HIV	ID	Weekend	HIV	ID	Weekend
1	A	E	E	A	E	H
2	A	E	H	A	E	A
3	B	F	F	B	H	G
4	B	F	H	B	H	I
5	A	G	A	A	G	F
6	A	G	E	A	G	C
7	B	C	C	A	F	B
8	B	C	G	D	C	G
9	D	I	D	B	C	D
10	D	I	H	B	D	H
11	A	B	B	A	B	F
12	A	B	I	A	B	A
13	C	F	F	C	H	H
14	C	F	I	C	H	I
15	A	H	H	B	I	C
16	A	H	D	B	I	E
17	B	E	B	A	E	D
18	B	E	H	A	E	E
19	A	I	I	A	C	F
20	A	I	D	A	C	C
21	C	D	C	B	G	A
22	C	D	I	B	G	C
23	B	F	F	C	F	D
24	B	F	G	C	F	C
25	A	H	A	C	C	G
26	A	H	H	D	I	D
27	C	D	D	A	B	E
28	C	D	E	A	B	I
29	A	B	A	B	D	D
30	A	B	B	B	D	A
31	C	E	C	C	F	E
32	C	E	A	C	F	F
33	B	D	D	B	F	I
34	B	D	E	B	I	C
35	A	I	A	A	G	G
36	A	I	G	A	G	I
37	D	C	C	D	C	A
38	D	C	F	D	D	E
39	B	E	E	A	B	D
40	B	E	B	B	I	I
41	A	G	G	B	I	G
42	A	G	E	D	F	F
43	C	I	C	C	F	C
44	C	I	I	D	E	I
45	A	F	A	D	E	E
46	A	F	F	A	B	A
47	B	G	B	A	B	D
48	B	G	I	A	D	G
49	D	C	D	A	D	F
50	D	C	B	B	G	G
51	B	G	G	B	G	E

The “HIV/ID” columns represent the assignments for the two concurrent services offered at the department. The “Weekend” column represents the assignments for weekend coverage in both services. Different colours and letters are used to distinguish different clinicians.

Table 6: Comparison of constraint satisfaction and objectives values in LP-generated and historical schedules.

	2015		2016		2017		2018	
	LP	Historical	LP	Historical	LP	Historical	LP	Historical
Constraint								
Block Coverage	✓	✓	✓	✓	✓	✓	✓	✓
Weekend Coverage	✓	✓	✓	✓	✓	✓	✓	✓
Min/Max	✓	✓	✓	✓	✓	✓	✓	✓
No Consecutive Blocks	✓		✓		✓		✓	
No Consecutive Weekends	✓		✓		✓	✓	✓	✓
Equal Weekends	✓		✓		✓		✓	
Equal Holidays	✓		✓	✓	✓		✓	✓
Objective								
Satisfied Block Requests	123/129	121/129	120/126	116/126	99/99	95/99	124/128	121/128
Satisfied Weekend Requests	113/113	111/113	119/119	112/119	75/75	75/75	115/115	113/115
Adjacent Block-Weekend Assignments	26/26	9/26	26/26	6/26	26/26	7/26	26/26	5/26

The objective denominator represents: the total number of requests submitted by clinicians for request objectives; and the total number of possible adjacencies for adjacency objectives.

3 concurrent services. However, when removing the NCB constraint, we saw a great improvement in runtime for divisions with 2 and 3 services, and we were able to generate a schedule with upwards of 50 clinicians in under 1.5 seconds.

Table 7: Comparison of program runtime (in seconds) for different numbers of services and total clinicians in the division.

	Number of Services				
Number of Clinicians	1	2	2 (*)	3	3 (*)
10	0.16	0.74	0.16	1.40	0.23
20	0.25	7468.86	0.32	–	0.43
30	0.42	–	0.49	–	0.66
50	0.62	–	0.82	–	1.14

Notes: “–” indicates that no solution was found within 24 hours; “(*)” indicates that the No Consecutive Blocks (NCB) constraint was removed.

For the remaining experiments, we simulated a department with 10 clinicians offering two services, similar to the department at St. Michael’s Hospital. The effect of an increasing number of requests per clinician on the runtime of the ILP solver is shown in Figure 1. In this experiment, each clinician was configured with 1 to 15 total block requests. The runtime of the algorithm is constant with respect to the number of requests, indicating that it can accommodate a lot of flexibility in clinician requests. Moreover, we see that all runs were completed in under 2 seconds.


Figure 2 presents the change in runtime when increasing the number of 2-week blocks in a department with 10 clinicians offering two services. In this experiment, we investigated time horizons from 5 to 110 blocks. This is equivalent to generating a schedule for up to 4 years ahead. The trend in the graph indicates a linear growth in runtime with respect to the time-horizon. Notably, the ILP solver was able to find all solutions in under 6 seconds, indicating very good performance for long-term scheduling.

6 Practical Considerations

In this section we discuss the lessons learned from developing the scheduling tool, as well as the implications of our approach as deployed in the hospital. The development of the scheduling solution started from a relaxed version of the problem (without constraints on consecutive blocks/weekends). We first investigated a maximum flow formulation for solving the problem. In this formulation, the problem is represented as a graph $G = (V, E)$ where the vertices comprise clinicians and blocks/weekends, and the edges comprise potential assignments of clinician to a block/weekend. A similar approach using maximum flow has been used in [29, 30, 11, 31]. The advantages of a maximum flow formulation are in the computational efficiency of the algorithm and the interpretability of the constraints and solution in terms of the graph [???]. However, a max flow formulation makes it difficult to incorporate new constraints. For instance, in order to accommodate the NCB and NCW constraints, additional edges and nodes need to be added to the graph, potentially reducing the computational efficiency and interpretability. As a result, we opted to reformulate the problem as an ILP, in which the additional constraints can be incorporated more easily.

The ILP formulation went through several stages of design and tweaking. For instance, we originally aimed to optimize only the BR and WR constraints. However, after receiving feedback on some of the test schedules, we learned that clinicians would prefer to have some of their weekend assignments closer to the block assignments. Once we were satisfied with the schedules produced by the algorithm, we targeted the user interface. Our goal was to develop an easy-to-use interface for the administrative staff to use, without requiring knowledge of the algorithm. [DL - do we need to elaborate more on this?]

Figure 1: Runtime of ILP solver with an increasing number of requests per clinician




Since the development of the tool, it has been used at St. Michael’s Hospital to generate schedules for the years 2019-21. The tool is used by administrative staff in the division in order to generate a potential schedule for the ID and HIV departments. The staff provides the following data as input to the scheduler: (1) maximum and minimum required number of blocks for each clinician, for each division; (2) a set of time-off requests for each clinician; (3) a set of holiday dates for the current calendar year. The supplied time-off requests in (2) are mapped to an appropriate block time-off request by the scheduler. For example, if a clinician requests time-off during weeks 2 and 3 of the year, the scheduler will interpret the request as time-off during block 1 (covering weeks 1-2) and 2 (covering weeks 3-4). The scheduler allows the user to disable certain constraints when generating a schedule in order to allow a broader range of solutions and prevent issues with infeasibility. It also allows the user to generate multiple schedules by randomly shuffling the input list of clinicians. Once a candidate schedule has been created, the department head may manually adjust the schedule to better accommodate some of the time-off requests and any additional requirements (although at the cost of breaking some of the hard constraints). Table 8 presents a comparison of the generated schedule and the final schedule after manual adjustments for the year 2019. We see that the manual adjustments included splitting or merging of certain blocks and local swaps between clinicians. Additionally, blocks and weekends may be adjusted manually due to newly hired clinicians or fellows being assigned to cover some of the assignments. For the year 2020, the generated schedules provided a good starting point for the schedule. However, due to the COVID-19 pandemic and an increase in the demand and workload of ID clinicians in the department, the schedules had to be adjusted manually as new requirements emerged. This is something that could be improved in future iterations of the model – specifically, the ability to (re-)generate shorter term schedules which are appropriately tuned based on a partial completed schedule provided as input. Overall, the generated schedule from the ILP model provide a good global template to which local adjustments can be easily performed post-hoc.

Table 8: Comparison of the automatically generated (ILP solution) schedule and the final schedule after manual adjustments for 2019.

Week #	ILP Solution			Manual Adjustment		
	HIV	ID	Weekend	HIV	ID	Weekend
1	D	C	C	D	E	E
2	D	C	D	D	E	F
3	L	I	L	D	L	D
4	L	I	F	D	L	K
5	C	E	E	C	I	H
6	C	E	G	C	I	G
7	D	L	D	D	E	L
8	D	L	F	D	E	L
9	I	K	K	C	I	F
10	I	K	I	C	I	I
11	D	C	C	D	K	D
12	D	C	G	D	K	G
13	L	K	K	C	G	H
14	L	K	E	C	G	E
15	C	F	C	L	L	K
16	C	F	F	L	L	J
17	I	L	L	I	C	L
18	I	L	I	I	C	E
19	D	H	D	D	K	I
20	D	H	H	D	K	D
21	C	F	F	I	F	G
22	C	F	E	I	F	E
23	D	I	I	D	C	I
24	D	I	H	D	C	D
25	C	E	C	I	F	H
26	C	E	E	I	F	G
27	I	G	G	D	E	E
28	I	G	I	D	E	F
29	D	H	H	C	H	K
30	D	H	K	C	H	J
31	C	L	L	D	K	H
32	C	L	G	D	K	D
33	I	K	I	C	C	I
34	I	K	F	L	C	K
35	D	C	C	I	H	I
36	D	C	D	D	H	F
37	L	G	G	L	K	D
38	L	G	K	I	K	H
39	D	E	E	L	G	G
40	D	E	L	L	G	I
41	C	G	G	L	H	K
42	C	G	H	I	H	E
43	L	F	F	C	F	E
44	L	F	D	I	F	F
45	D	K	K	D	G	G
46	D	K	I	L	G	L
47	C	L	L	C	K	D
48	C	L	D	C	K	H
49	D	H	H	D	I	L
50	D	H	K	D	L	K
51	C	K	C	C	L	H

Figure 2: Runtime of ILP solver with an increasing number of 2-week blocks



7 Discussion

In this paper, we present a simple, yet flexible, integer linear programming formulation to generate schedules for clinical departments at hospitals. The challenge in applying ILP to the task of scheduling clinicians lies in the computational complexity of finding an optimal solution. As the size of the scheduling problem grows, due to a larger roster of clinicians or more complicated constraints, the time it takes to generate an optimal schedule may grow exponentially in the worst case. Many previous approaches to creating schedules in similar scenarios have avoided this problem by using heuristics to find an approximately optimal solution in a shorter amount of time [18].

We presented a formulation that includes both hard constraints to ensure the schedule satisfies hospital and logistics requirements, and a multi-goal objective function to satisfy soft constraints (work preferences of clinicians). Although we restricted our application of the formulation to a set of constraints for the particular needs of the case study (St. Michael’s Hospital Division of Infectious Diseases), our formulation can be adapted to various clinical departments at different hospitals. The flexibility of our ILP allows changing the number of services provided in a division, the length of a work block, clinicians’ preference for block to weekend adjacency as well as clinicians’ requests for time off.

When comparing the optimal schedule generated by our tool to the manually-created schedules at St. Michael’s Hospital, we found that the ILP formulation was always able to find an optimal schedule satisfying all required hard constraints, unlike the manual schedule, which often did not satisfy all constraints. Moreover, due to the multi-goal objective function in the ILP, the algorithm was able to fulfill the majority of clinician preferences and requests, more so than the manually-created schedule. These observations reinforce the benefits of automated tools when generating schedules in hospital departments to balance the workload of clinicians and improve the service provided to patients. The use of automated tools alleviates the time spent on designing the schedule by hand, and provides clinical departments with a more fair distribution of work that helps improve the overall satisfaction of both employees and patients [32].

In our simulations, we also found that increasing the number of requests per clinician did not affect the

runtime of the algorithm, highlighting the flexibility of the tool to incorporate clinician preferences. Further, we saw that the algorithm can accommodate an increase in time-horizon up to four years with little impact on runtime, suggesting the algorithm can be used to generate schedules far in advance. A key limitation we identified was the sensitivity of the runtime to larger numbers of services offered by a single department. One solution to mitigate the runtime issues created by a larger number of services would be removing the constraint that prevents assignment of consecutive blocks, followed by manual readjustment from the generated schedule. Overall, our sensitivity analyses using simulated data provided reassurance that the ILP formulation can be applied to schedule clinicians across real-world variability between clinical departments. Next steps include expanding the generalizability of the tool beyond smaller clinical departments to larger departments within and outside of health-care – especially those that provide multiple services in parallel for patients and other clients. As well, additional work can be done to incorporate other clinician preferences, such as the ability to request time-on slots or preference for time of year.

8 Acknowledgements

The authors are grateful to Julie Veitch (St. Michael’s Hospital) for her contributions to testing and designing the scheduling software. SM is supported by a Canadian Institutes of Health Research New Investigator Award. JK is supported by a Natural Sciences and Engineering Research Council doctoral award. DL conducted part of this project as a Keenan Research Summer Student, Li Ka Shing Knowledge Institute, St. Michael’s Hospital, University of Toronto.

9 Funding

The work was jointly funded by the Ontario Ministry of Science and Innovation Early Researcher Award Number ER17-13-043; the Division of Infectious Diseases, St. Michael’s Hospital, University of Toronto; and the University of Toronto Work Study Program.

10 Conflict of Interest

On behalf of all authors, the corresponding author states that there is no conflict of interest.

References

- [1] Uwe Aickelin, Edmund K. Burke, and Jingpeng Li. Improved Squeaky Wheel Optimisation for Driver Scheduling. en. In Thomas Philip Runarsson, Hans-Georg Beyer, Edmund Burke, Juan J. Merelo-Guervós, L. Darrell Whitley, and Xin Yao, editors, *Parallel Problem Solving from Nature - PPSN IX*, Lecture Notes in Computer Science, pages 182–191. Springer Berlin Heidelberg, 2006. ISBN: 978-3-540-38991-0.
- [2] Asvin Goel, Claudia Archetti, and Martin Savelsbergh. Truck driver scheduling in Australia. *Computers & Operations Research*, 39(5):1122–1132, May 2012. ISSN: 0305-0548. DOI: 10.1016/j.cor.2011.05.021.
- [3] Maik Günther and Volker Nissen. Combined Working Time Model Generation and Personnel Scheduling. en. In Wilhelm Dangelmaier, Alexander Blecken, Robin Delius, and Stefan Klöpfer, editors, *Advanced Manufacturing and Sustainable Logistics*, Lecture Notes in Business Information Processing, pages 210–221. Springer Berlin Heidelberg, 2010. ISBN: 978-3-642-12494-5.

- [4] Nicolas Chapados, Marc Joliveau, and Louis-Martin Rousseau. Retail Store Workforce Scheduling by Expected Operating Income Maximization. en. In Tobias Achterberg and J. Christopher Beck, editors, *Integration of AI and OR Techniques in Constraint Programming for Combinatorial Optimization Problems*, Lecture Notes in Computer Science, pages 53–58. Springer Berlin Heidelberg, 2011. ISBN: 978-3-642-21311-3.
- [5] Volker Nissen and Maik Günther. Automatic Generation of Optimised Working Time Models in Personnel Planning. en. In Marco Dorigo, Mauro Birattari, Gianni A. Di Caro, René Doursat, Andries P. Engelbrecht, Dario Floreano, Luca Maria Gambardella, Roderich Groß, Erol Şahin, Hiroki Sayama, and Thomas Stützle, editors, *Swarm Intelligence*, Lecture Notes in Computer Science, pages 384–391. Springer Berlin Heidelberg, 2010. ISBN: 978-3-642-15461-4.
- [6] M. E. T. Horn, H. Jiang, and P. Kilby. Scheduling patrol boats and crews for the Royal Australian Navy. *Journal of the Operational Research Society*, 58(10):1284–1293, October 2007. ISSN: 0160-5682. DOI: 10.1057/palgrave.jors.2602300.
- [7] Manuel Laguna and Terry Wubben. Modeling and Solving a Selection and Assignment Problem. en. In Bruce Golden, S. Raghavan, and Edward Wasil, editors, *The Next Wave in Computing, Optimization, and Decision Technologies*, Operations Research/Computer Science Interfaces Series, pages 149–162. Springer US, 2005. ISBN: 978-0-387-23529-5.
- [8] M.N. Azaiez and S.S. Al Sharif. A 0-1 goal programming model for nurse scheduling. en. *Computers & Operations Research*, 32(3):491–507, March 2005. ISSN: 03050548. DOI: 10.1016/S0305-0548(03)00249-1.
- [9] Lorraine Trilling, Alain Guinet, and Dominiue Le Magny. Nurse scheduling using integer linear programming and constraint programming. *IFAC Proceedings Volumes*, 39(3):671–676, 2006.
- [10] Maya Widyastiti, Amril Aman, and Toni Bakhtiar. Nurses Scheduling by Considering the Qualification using Integer Linear Programming. en. *TELKOMNIKA (Telecommunication Computing Electronics and Control)*, 14(3):933, September 2016. ISSN: 2302-9293, 1693-6930. DOI: 10.12928/telkomnika.v14i3.2913.
- [11] Ahmed Ali El Adoly, Mohamed Gheith, and M. Nashat Fors. A new formulation and solution for the nurse scheduling problem: A case study in Egypt. en. *Alexandria Engineering Journal*, 57(4):2289–2298, December 2018. ISSN: 11100168. DOI: 10.1016/j.aej.2017.09.007.
- [12] Uwe Aickelin and Kathryn A. Dowland. Exploiting problem structure in a genetic algorithm approach to a nurse rostering problem. en. *Journal of Scheduling*, 3(3):139–153, 2000. ISSN: 1099-1425. DOI: 10.1002/(SICI)1099-1425(200005/06)3:3<139::AID-JOS41>3.0.CO;2-2.
- [13] A. Jan, M. Yamamoto, and A. Ohuchi. Evolutionary algorithms for nurse scheduling problem. In *Proceedings of the 2000 Congress on Evolutionary Computation. CEC00 (Cat. No.00TH8512)*, volume 1, 196–203 vol.1, July 2000. DOI: 10.1109/CEC.2000.870295.
- [14] H. Kawanaka, K. Yamamoto, T. Yoshikawa, T. Shinogi, and S. Tsuruoka. Genetic algorithm with the constraints for nurse scheduling problem. In *Proceedings of the 2001 Congress on Evolutionary Computation (IEEE Cat. No.01TH8546)*, volume 2, 1123–1130 vol. 2, May 2001. DOI: 10.1109/CEC.2001.934317.
- [15] Andrzej Jaskiewicz. A metaheuristic approach to multiple objective nurse scheduling. *Foundations of Computing and Decision Sciences*, 22(3):169–184, 1997.
- [16] Slim Abdennadher and Hans Schlenker. Nurse scheduling using constraint logic programming. In *AAAI/IAAI*, pages 838–843, 1999.
- [17] Haibing Li, Andrew Lim, and Brian Rodrigues. A Hybrid AI Approach for Nurse Rostering Problem. In *Proceedings of the 2003 ACM Symposium on Applied Computing, SAC '03*, pages 730–735, New York, NY, USA. ACM, 2003. ISBN: 978-1-58113-624-1. DOI: 10.1145/952532.952675.

- [18] Edmund K. Burke, Patrick De Causmaecker, Greet Vanden Berghe, and Hendrik Van Landeghem. The State of the Art of Nurse Rostering. en. *Journal of Scheduling*, 7(6):441–499, November 2004. ISSN: 1094-6136. DOI: 10.1023/B:JOSH.0000046076.75950.0b.
- [19] Tim B. Cooper and Jeffrey H. Kingston. The complexity of timetable construction problems. en. In Gerhard Goos, Juris Hartmanis, Jan Leeuwen, Edmund Burke, and Peter Ross, editors, *Practice and Theory of Automated Timetabling*. Volume 1153, pages 281–295. Springer Berlin Heidelberg, Berlin, Heidelberg, 1996. ISBN: 978-3-540-61794-5 978-3-540-70682-3. DOI: 10.1007/3-540-61794-9_66.
- [20] Patrick De Causmaecker and Greet Vanden Berghe. A categorisation of nurse rostering problems. en. *Journal of Scheduling*, 14(1):3–16, February 2011. ISSN: 1099-1425. DOI: 10.1007/s10951-010-0211-z.
- [21] P. L. Hammer and S. Rudeanu. *Boolean Methods in Operations Research and Related Areas*. en. Ökonometrie und Unternehmensforschung Econometrics and Operations Research. Springer-Verlag, Berlin Heidelberg, 1968. ISBN: 978-3-642-85825-3.
- [22] Ivan P. Stanimirovic, Milan Lj Zlatanovic, and Marko D. Petkovic. On the linear weighted sum method for multi-objective optimization. *Facta Acta Univ*, 26(4):49–63, 2011.
- [23] John E. Mitchell. Branch-and-cut algorithms for combinatorial optimization problems. *Handbook of applied optimization*, 1:65–77, 2002.
- [24] Ron Shamir. The Efficiency of the Simplex Method: A Survey. *Management Science*, 33(3):301–334, 1987. ISSN: 0025-1909.
- [25] Burak Bilgin, Patrick De Causmaecker, Benoît Rossie, and Greet Vanden Berghe. Local search neighbourhoods for dealing with a novel nurse rostering model. en. *Annals of Operations Research*, 194(1):33–57, April 2012. ISSN: 1572-9338. DOI: 10.1007/s10479-010-0804-0.
- [26] Edmund K. Burke, Patrick De Causmaecker, Sanja Petrovic, and Greet Vanden Berghe. Metaheuristics for Handling Time Interval Coverage Constraints in Nurse Scheduling. *Applied Artificial Intelligence*, 20(9):743–766, December 2006. ISSN: 0883-9514. DOI: 10.1080/08839510600903841. Publisher: Taylor & Francis _eprint: <https://doi.org/10.1080/08839510600903841>.
- [27] David Landsman. On-call scheduling tool for clinicians. <https://github.com/c-uhs/scheduler>, 2019.
- [28] johnjforrest, Stefan Vigerske, Ted Ralphs, Haroldo G. Santos, Lou Hafer, Bjarni Kristjansson, jpfasano, Edwin Straver, Miles Lubin, rlougee, jgoncal1, h-i-gassmann, and Matthew Saltzman. Coin-or/Cbc: Version 2.9.9, 2019. DOI: 10.5281/zenodo.3246628.
- [29] Peter Brucker, Rong Qu, and Edmund Burke. Personnel scheduling: models and complexity. *European Journal of Operational Research*, 210(3):467–473, 2011. Publisher: Elsevier.
- [30] Pieter Smet, Peter Brucker, Patrick De Causmaecker, and Greet Vanden Berghe. Polynomially solvable personnel rostering problems. *European Journal of Operational Research*, 249(1):67–75, 2016. Publisher: Elsevier.
- [31] Margarida Moz and Margarida Vaz Pato. An integer multicommodity flow model applied to the rerostering of nurse schedules. *Annals of Operations Research*, 119(1):285–301, 2003. Publisher: Springer.
- [32] Rhian Silvestro and Claudio Silvestro. An evaluation of nurse rostering practices in the National Health Service. en. *Journal of Advanced Nursing*, 32(3):525–535, 2000. ISSN: 1365-2648. DOI: 10.1046/j.1365-2648.2000.01512.x.